http://cs.nyu.edu/courses/summer15/CSCI-GA.2250-001/labs/lab1/

You are to implement a **two-pass linker** and submit the **source** code, which we will compile and run. Email your source code to the assigned TA with instructions on how to build your program (preferably as a Makefile). **Please do not submit inputs or outputs.** Your program must take one input parameter which will be the name of an input file to be processed. All output should go to standard output. The languages of choice for this first lab are C/C++/Java. All subsequent labs will be C/C++ only. You may develop your lab on any machine you wish, but you must ensure that it compiles and runs on the NYU system assigned to the course (energon1/2) where it will be graded. It is your responsibility to make sure it executes on those machines. Note, when you work on energon1 or energon2 the default GCC/G++ compiler is v4.4.7. If you use advanced features there is a version 4.6 and 4.8; use gcc46, gcc48, g++46 or g++48 instead.

In general, a linker takes individually compiled code/object modules and creates a single executable by resolving external symbol references (e.g. variables and functions) and module relative addressing by assigning global addresses after placing the modules' object code at global addresses.

We assume a target machine with the following properties: (a) word addressable, (b) addressable memory of 512 words, and (c) each word consisting of 4 decimal digits. [*I know that is a really strange machine*].

The input to the linker is a file containing a **sequence of tokens** (symbols and integers and instruction type characters). Don't assume tokens that make up a section to be on one line, don't make assumptions about how much space separates tokens or that lines are non-empty for that matter or that each input conforms syntactically. Symbols always begin with alpha characters followed by optional alphanumerical characters, i.e.[a-Z][a-Z0-9]*. Symbols can be up to 16 characters. Integers are decimal based. Instruction type characters are (I, A, R, E). Token delimiters are ' ', '\t' or '\n'.

The input file to the linker is structured as a series of "object module" definitions.
Each "object module" definition contains three parts (in fixed order): definition list, use list, and program text.

- *definition list* consists of a count *defcount* followed by *defcount* pairs (S, R) where S is the symbol being defined and R is the relative word address (offset) to which the symbol refers in the module.
- *use list* consists of a count *usecount* followed by *usecount* symbols that are referred to in this module. These could include symbols defined in the *definition list* of any module (prior or subsequent or not at all).
- *program text* consists of a count *codecount* followed by *codecount* pairs (**type, instr**), where *instr* is a 4-digit instruction (integer) and *type* is a single character indicating **I**mmediate, **A**bsolute, **R**elative, or **E**xternal. *codecount* is thus the length of the module.

An instruction (upto 4 decimals digits) is composed of an opcode (leftmost digit) and an operand (rightmost 3 digits). The opcode always remains unchanged by the linker.
The operand is modified/retained based on the instruction type in the *program text* as follows:
(**I**) an immediate operand is unchanged; Note that there is no opcode in this case and operand is all four digits
(**A**) an absolute address is unchanged;
(**R**) a relative address is relocated by replacing the relative address with the absolute address of that relative address after the modules global address has been determined.
(**E**) an external address is an index into the uselist. For example, a reference in the program text with operand K represents the Kth symbol in the use list, using 0-based counting, e.g., if the use list is ''2 f g'', then an instruction ''E 7000'' refers to f, and an instruction ''E 5001'' refers to g. You must identify to which global address the symbol is assigned and then replace the address.

The linker must process the input twice (that is why it is called two-pass) (to preempt the favored question: "Can I do it in one pass?" → NO). **Pass One** parses the input and verifies the correct syntax and determines the base address for each module and the absolute address for each defined symbol, storing the latter in a symbol table. The first module has base address zero; the base address for module X+1 is equal to the base address of module X plus the length of module X. The absolute address for symbol S defined in module M is the base address of M plus the relative address of S within M. After pass one print the symbol table (including errors related to it (see rule2 later)). Do not store parsed tokens, only store meta data (e.g. deflist, uselist, num-instructions)

Pass Two again parses the input and uses the base addresses and the symbol table entries created in pass one to generate the actual output by relocating relative addresses and resolving external references.
You must clearly mark your two passes in the code through comments and/or proper function naming.

Other requirements: error detection, limits, and space used.

To receive full credit, you must check the input for various errors. All errors/warnings should follow the message catalog provided below. We will do a textual difference against a reference implementation to grade your program. Any reported difference will indicate a non-compliance with the instructions provided and is reported as an error and result in deductions.
You should continue processing after encountering an error/warning (other than a syntax error) and you should be able to detect multiple errors in the same run.

1. You should stop processing if a syntax error is detected in the input, print a syntax error message with the line number and the character offset in the input file where observed. A syntax error is defined as a missing token (e.g. 4 used symbols are defined but only 3 are given) or an unexpected token. Stop processing and exit.
2. If a symbol is defined multiple times, print an error message and use the value given in the first definition. Error message to appear as part of printing the symbol table (following symbol=value printout on the same line)
3. If a symbol is used in an E-instruction but not defined, print an error message and use the value zero.
4. If a symbol is defined but not used, print a warning message and continue.
5. If an address appearing in a definition exceeds the size of the module, print an error message and treat the address given as 0 (relative to the module).
6. If an external address is too large to reference an entry in the use list, print an error message and treat the address as immediate.
7. If a symbol appears in a use list but it not actually used in the module (i.e., not referred to in an E-type address), print a warning message and continue.
8. If an absolute address exceeds the size of the machine, print an error message and use the absolute value zero.
9. If a relative address exceeds the size of the module, print an error message and use the module relative value zero (that means you still need to remap "0" that to the correct absolute address).
10. If an illegal immediate value (I) is encountered (i.e. more than 4 numerical digits), print an error and convert the value to 9999.
11. If an illegal opcode is encountered (i.e. more than 4 numerical digits), print an error and convert the <opcode,operand> to 9999.

The following exact limits are in place.
   a) Accepted symbols should be upto 16 characters long (not including terminations e.g. '\0'), any longer symbol names are erroneous.
   b) a uselist or deflist should support 16 definitions, but not more and an error should be raised.
   c) number instructions are unlimited (hence the two pass system), but in reality they are limited to the machine size.
   d) Symbol table should support at least 256 symbols

There are several sample inputs and outputs provided as part of the sample input files / output files (see website).
The first (*input-1)* is shown below and the second (*input-2)* is a re-formatted version of the first. They both produce the same output as the input is token-based and hence present the same content to the linker. Some of the input sets contain errors that you are to detect as described above. Note that when you have questions regarding error, please first make sure you structure of the input is not messing with your mind.
We will run your lab on these (and other) input sets. Please submit the SOURCE code for your lab, together with a README file (required) describing how to compile and run it. Your program must accept one command line argument giving the name of the input file (which must accept a full path as we are running this through a grading harness);

```
1 xy 2
2 z xy
5 R 1004  I 5678  E 2000  R 8002  E 7001
0
1 z
6 R 8001  E 1000  E 1000  E 3000  R 1002  A 1010
0
1 z
```

```
2 R 5001  E 4000
1 z 2
2 xy z
3 A 8000  E 1001  E 2000
```

Your output is expected to strictly follow this format (with exception of empty lines):

```
Symbol Table
xy=2
z=15

Memory Map
000: 1004
001: 5678
002: 2015
003: 8002
004: 7002
005: 8006
006: 1015
007: 1015
008: 3015
009: 1007
010: 1010
011: 5012
012: 4015
013: 8000
014: 1015
015: 2002
```

The following output that is heavily annotated for clarity and class discussion. Your output is **not** expected to be this fancy. It should help you understand the operation and mapping of symbols etc.

```
Symbol Table
xy=2
z=15
Memory Map
+0
0:      R 1004         1004+0 =  1004
1:      I 5678                   5678
2: xy:  E 2000 ->z               2015
3:      R 8002         8002+0 =  8002
4:      E 7001 ->xy              7002
+5
0:      R 8001         8001+5 =  8006
1:      E 1000 ->z               1015
2:      E 1000 ->z               1015
3:      E 3000 ->z               3015
4:      R 1002         1002+5 =  1007
5:      A 1010                   1010
+11
0:      R 5001         5001+11=  5012
1:      E 4000 ->z               4015
+13
0:      A 8000                   8000
1:      E 1001 ->z               1015
2 z:    E 2000 ->xy              2002
```

Note that even an empty program should have the "Symbol Table" and "Memory Map" line.

We grade by using a "diff –b –B -E" against the reference output created by my test program using a grading harness.
Inputs will be the ones provided on the web as well as other once that will be checked for several of the error conditions. It is imperative that you match the output as generated by the ref program to allow for automated testing.
For a test case to pass you must catch ALL warning/errors and generate the correct output.

Example: (note the

```
Symbol Table
X21=3
X31=4

Memory Map
000: 1003
001: 1003
002: 1003
003: 2000  Error: Absolute address exceeds machine size; zero used
004: 3000  Error: Relative address exceeds module size; zero used

Warning: Module 3: X31 was defined but never used
```

Parse error should abort processing.
Error messages must be following the instruction as shown above
Warnings should be printed at the end after the memory map and in order of module appearance (note modules are numbered starting with '1') or in the case of rule 5 print it before the Symbol Table is printed (see input-10).

I provide in C the code to print parse errors, which also gives you an indication what is considered a parse error.

```
void __parseerror(int errcode) {
    static char* errstr[] = {
        "NUM_EXPECTED",              // Number expect
        "SYM_EXPECTED",              // Symbol Expected
        "ADDR_EXPECTED",             // Addressing Expected
        "SYM_TOLONG",                // Symbol Name is to long
        "TO_MANY_DEF_IN_MODULE",     // > 16
        "TO_MANY_USE_IN_MODULE",     // > 16
        "TO_MANY_INSTR",             // total num_instr exceeds memory size (512)
    };
    printf("Parse Error line %d offset %d: %s\n", linenum, lineoffset, errstr[errcode]);
}
```

(Note: line numbers start with 1 and offsets in the line start with 1, offsets should indicate the first character offset of the token that is wrong, not the last). Tabs count as one character.
Error messages have the following text and should appear right at the end of the line you are printing out

"Error: Absolute address exceeds machine size; zero used"	(see rule 8)
"Error: Relative address exceeds module size; zero used"	(see rule 9)
"Error: External address exceeds length of uselist; treated as immediate"	(see rule 6)
"Error: %s is not defined; zero used" (insert the symbol name for %s)	(see rule 3)
"Error: This variable is multiple times defined; first value used"	(see rule 2)
"Error: Illegal immediate value; treated as 9999"	(see rule 10)
"Error: Illegal opcode; treated as 9999"	(see rule 11)

Warnings have the following text and are on a separate line at the end after the memory map printout. Exception is the warning for rule 5 which should be printed after each module has been parsed in pass-1 and hence should appear in order of appearance and before the symbol table is printed.

"Warning: Module %d: %s to big %d (max=%d) assume zero relative\n"	(see rule 5)
"Warning: Module %d: %s appeared in the uselist but was not actually used\n"	(see rule 7)
"Warning: Module %d: %s was defined but never used\n"	(see rule 4)

Parse Error Location:

Parse errors are to be located at the first character of the wrong token, or if end-of-file is reached at the end of file location.
There is one special case when the eof ends with `\n'. My expectation is that the line number reported actually exists in the file and that an editor (e.g. vi) can jump to it. In this particular case the linenumber to be reported is the last line read and the last position of that line, not the next line and offset 1. (see *input-12* for an example. The error is at the very last position of the line. Reason is when one does a linecount on the file ("wc –l input-12") it shows 3.

Grading and Testing your program.

As part of the labsamples.tar on the labs/lab1 website you will see a *runit.sh* and a *gradeit.sh* script (the same we will use for grading albeit with more inputs).

Execute as follows:
Create yourself a directory <your-outdir> where your outputs will be created.
> cd labsamples
> ./runit.sh <your-outdir> <your-executable and optional arguments>
The above will create all the outputs for the available inputs (1-19)
> ./gradeit.sh . <your-outdir>
The above will compare the reference output (out-[1-19]) with the one created with your program and tell you how many you got right.
There will be a file called <your-outdir>/LOG that contains which cases you got wrong and where the differences are.
If you want to analyze further please run the "diff –b –B –E" by hand on a particular output pair.

Please only submit your source code and instructions.

Writing a parser:

Here are some hints on writing a parser. In general one could write this parser using   lex/yacc, however this is so simple that I suggest writing a simple recursive decent parser, in particular since you have to parse the input twice. It would have a structure similar to this (all pseudoCode). This structure is simply copied twice and the actions taken in each path are different.

ReadFile() { while (!eof) { createModule() = { readDefList(); readUseList(); readInstList(); } }
ReadDefList() { numDefs=readInt(); for (i=0;i<numDefs(); i++) { readDef(); }
ReadDef() { str= readSym(); val=readInt(); createSymbol(str,val); }

Ofcourse you will have to deal with the errors etc.
The first pass creates the metadata and does error checking, the second pass performs additional error checking and instruction transformation.